Assessment of a Plain-Language House Style System Through an AI-Writing Detection Taxonomy

A Practical Academic-Style Evaluation for Executive, Educational, Social, and Informal Writing Use Cases

Prepared by **Kalen Howell Sr.** | www.kalenhowellsr.com

Prepared through research, synthesis, and review in collaboration with AI. Final judgment, framing, and publication responsibility remain with Kalen Howell Sr.

Prepared from the uploaded house style source pack and Russell, Karpinska, and Iyyer (2025)

Abstract

This assessment evaluates a compact plain-language house style system through the lens of a recent taxonomy of clues used by expert human detectors of AI-generated text. The system is already strong for business and executive writing because it rewards directness, short prose, defined terms, decision structure, and a light checkpoint workflow. Its main weakness is enforcement. The written guidance is broad and useful, while the automated Vale layer currently checks only a small portion of the rules. The second weakness is domain breadth. Student essays, social media posts, long-form reports, and informal writing need additional mode-specific checks because each domain carries different expectations for voice, evidence, structure, and risk. The recommendation is to keep the system small but add a few high-value controls: a compact domain-mode checklist, a safer evidence and attribution section, objective Vale warnings for sentence length, filler phrases, acronym candidates, and repetitive conclusions, plus a clear exception process. The system should not become an AI detector or a humanizer. It should become a quality and risk-control framework that reduces generic, formulaic, over-polished writing while preserving speed and author judgment.

Keywords: plain language, house style, AI-generated text, prose linting, executive writing, student essays, social media, business writing, long-form reports

1. Introduction

The central question is whether your house style system can do more than make prose clear. The deeper question is whether it can help writers avoid the patterns that make prose feel generic, over-processed, or machine-like across several writing domains. My assessment is that it can, but only if the taxonomy from Russell, Karpinska, and Iyyer is used as a diagnostic lens rather than a rigid rulebook.

The uploaded system has a disciplined design. It separates core house style from executive writing guidance and uses a style gate only at deliverable checkpoints. That architecture protects delivery speed and lowers adoption friction. The current automation is intentionally light: Vale checks selected deliverable Markdown and text files, warns when Vale is missing by default, and applies a small set of local rules for jargon substitutions, Unicode em dashes, and one contrast formula pattern.

The AI-writing paper adds a new lens. Its taxonomy shows that expert human detectors do not rely only on obvious vocabulary. They also notice sentence patterns, quotes, over-polished grammar, lack of originality, unnatural clarity, repetitive formatting, formulaic introductions and conclusions, tone, names and titles, factuality, and topic avoidance. In the study, five expert annotators, defined by frequent use of large language models for writing tasks, misclassified only one of 300 nonfiction articles by majority vote. That result is impressive, but it is bounded. The paper studied short American English nonfiction articles, not student essays, social posts, business deliverables, long technical reports, or informal messages.

The practical goal, therefore, is not to build a detector. The practical goal is to strengthen your style system so it catches predictable prose weaknesses, improves authenticity and decision quality, and reduces false confidence. The final design should remain lightweight. It should help writers revise the work they already produced rather than turn writing into a compliance exercise.

2. Evidence Base and Method

This assessment uses four source groups. First, it uses the uploaded house style source pack: the system summary, the repository and writing guides, the style gate and Vale rules, and the analysis brief. Second, it uses the AI-writing detection paper by Russell, Karpinska, and Iyyer. Third, it uses public plain-language and usability guidance, including ISO 24495-1, Digital.gov, the U.S. Department of Labor quick reference guide, the Microsoft Writing Style Guide, and Nielsen Norman Group research on concise and scannable writing. Fourth, it uses current AI-detection robustness research, especially RAID and Sadasivan et al., to avoid overstating what any style or detector can prove.

The rating scale is intentionally simple. A 1.0 means the category is absent, actively risky, or unsupported. A 5.0 means the category is present but uneven, mostly manual, or domain-limited. A 10.0 means the category is clear, practical, consistently applied, supported by the workflow, and safe across intended uses. Scores reflect the current system, not the authorial skill of a specific writer.

The score also separates guidance strength from enforcement strength. This matters because a style guide can be well written while the gate enforces little of it. Your system is in that position today. The guidance is much stronger than the automation. That is acceptable if authors are trained and review practices are consistent. It is risky if Codex, Claude, and human contributors all produce deliverables without a shared review checklist.

3. Current System Map

The current system has three operating layers. The core layer is the plain-language house style: short sentences, short paragraphs, plain words, defined acronyms, no Unicode em dashes, and no contrast formulas that bury the point. The executive layer adds decision structure: lead with the answer, explain what changed or what matters, show risks and tradeoffs, make the recommendation easy to find, and state next steps. The automation layer runs Vale at deliverable checkpoints and intentionally leaves rough internal work alone.

This architecture is sound. It fits a docs-first repository and a mixed authoring environment where agents and humans both contribute. It also aligns with external plain-language principles. ISO 24495-1 frames plain language around

whether readers can find, understand, and use the information. The Department of Labor guide similarly emphasizes audience purpose, leading with the main point, short sentences and paragraphs, active voice, limited acronyms, and review. The Microsoft guide supports simple sentence structures and consistent terminology.

The system's most important design choice is checkpoint-based enforcement. Style checks happen after editing deliverable prose and before final handoff. That preserves momentum during drafting and prevents rough notes from becoming over-managed. The tradeoff is that late-stage checking can miss structural problems that would have been easier to correct during planning. For executive deliverables, this can be mitigated with a short pre-draft checklist rather than more lint rules.

System Layer	Current Function	Strength	Primary Gap
Core house style	Defines plain language rules for all human-facing prose.	Short, memorable, and broadly applicable.	Some rules are not measurable or not enforced.
Executive guide	Adds structure for memos, decision briefs, recommendations, and updates.	Strong decision orientation and clear anti-patterns.	Manual judgment is required for recommendation quality and tradeoffs.
Style gate script	Runs Vale against selected deliverable files at checkpoints.	Light, scoped, and low-friction.	Soft default can hide missing automation.
Vale rules	Enforces em dashes, limited jargon substitutions, and one contrast formula pattern.	Objective checks are easy to maintain.	Coverage is narrow relative to the written rules.
Operating model	Allows rough notes to stay rough and gates final prose.	Protects speed and adoption.	Needs a clearer exception and review policy.

4. The AI-Writing Detection Taxonomy as a Style Risk Lens

The taxonomy in the Russell et al. paper is useful because it names the features that make AI-generated prose visible to experienced readers. The top categories are vocabulary, sentence structure, grammar and punctuation, originality, quotes, and clarity, followed by formatting, conclusions, formality, names and titles, tone, introductions, factuality, topics, and miscellaneous cues. The most common category was vocabulary, but the paper is more valuable because it goes beyond vocabulary.

The main strategic insight is that AI-like prose is often too balanced, too tidy, too frictionless, and too generically positive. It may use the same polished quote structure, the same optimistic ending, the same safe vocabulary, and the same neatly symmetrical sentence patterns. Those traits can be harmless in low-stakes writing, but they weaken trust in high-stakes writing because readers cannot see the author's judgment, evidence, or lived understanding of the issue.

Your house style already resists several of these signals. It discourages jargon, filler, long paragraphs, contrast formulas, and vague executive summaries. It requires the main point early and asks writers to separate what is known, inferred, and recommended. Those rules directly push against generic AI prose. The system is weaker where the taxonomy requires domain judgment: originality, quotes, named entities, factual texture, tone variation, informal voice, and document-specific conventions.

Taxonomy Cue	Current Coverage	Score	Practical Enhancement
Vocabulary	Plain language rule and small substitution list	7.0	Expand carefully. Flag filler and buzzwords as suggestions, not errors.
Sentence Structure	Short direct sentences are stated, but not enforced	6.2	Add sentence-length and comma-count warnings. Keep as suggestions.
Grammar and Punctuation	Em dash ban is enforced; grammar quality is manual	5.4	Keep hard rule for em dashes. Avoid enforcing false grammar

			preferences.
Originality	Not addressed directly	4.4	Add a domain prompt for specific examples, judgment, and non-generic insight.
Quotes	Not addressed	2.5	For essays and reports, require quote purpose, attribution, and integration.
Clarity	Central design goal	8.4	Add examples showing clear but not bland prose.
Formatting	Repository rules and executive structure provide partial coverage	6.8	Add simple templates for reports and decision briefs.
Conclusions	Executive guide avoids generic summaries	8.0	Add a lint warning for empty wrap-up phrases.
Formality	Professional tone implied; informal mode absent	6.0	Create a channel-specific voice note for informal writing.
Names and Titles	Specificity encouraged through numbers, owners, dates, and examples	6.3	Add source-check expectations for named people, products, and claims.
Tone	Executive voice is clear and disciplined	7.3	Allow more human variance in essays, social posts, and informal notes.
Introductions	Strong rule: lead with signal	8.7	Keep. Add domain alternatives for essays and social posts.
Factuality	Known, inferred, recommended separation is strong	7.6	Add evidence ledger for reports and academic work.
Topics	Risks and tradeoffs are required in executive work	7.1	Ensure uncomfortable facts are not smoothed away.

5. Overall Rating of the Current House Style System

The current system earns an overall rating of 7.3 out of 10.0. That score is good. It reflects a strong, practical writing philosophy and a mature understanding of workflow friction. It also reflects a gap between the written system and the automated system. The highest leverage improvement is not a larger style guide. It is better coverage of a few objective checks and a small set of domain modes.

Category	Score	Band	Assessment
Purpose and scope	8.8	Strong	The system clearly protects comprehension, decision quality, and delivery speed.
Core plain-language rules	8.2	Strong	The rules are short, memorable, and aligned with plain-language guidance.
Executive writing structure	8.9	Strong	The structure supports leadership decisions and avoids weak recommendations.
AI-signature risk coverage	6.8	Developing	The system resists generic prose but does not yet address voice, quotes, originality, or

			factual texture.
Domain adaptability	6.3	Developing	Business writing is well supported. Other domains need mode-specific additions.
Evidence and attribution discipline	7.2	Good	Known, inferred, and recommended separation is strong. Citation and source handling need clearer rules.
Automation coverage	4.7	Weak	Vale enforces a small subset of the written guidance.
Workflow usability	8.6	Strong	The checkpoint model is light and practical.
Exception handling	5.8	Developing	Intentional acceptance of warnings is implied but not formalized.
Training and examples	4.9	Weak	The system needs before-and-after examples to make judgment rules easier to learn.

6. Domain-Specific Application

The five requested domains require different interpretations of the same core style. A single rule set will either overfit business writing or flatten more personal domains. The right solution is a small mode layer: each domain keeps the core house style but adds three to five domain-specific checks. This gives practical strength without creating a heavy governance model.

Domain	Current Score	Practical Target	Reason
Student essays	6.1	7.7	Strong for clarity, weak for academic voice, evidence, citation practice, and student authorship risk.
Social media posts	5.8	7.4	Strong for brevity, weak for platform voice, hooks, engagement context, and authenticity.
Business documents	8.6	9.2	Best fit. The executive guide maps directly to business decisions.
Long-form reports	7.2	8.5	Good foundation. Needs stronger hierarchy, evidence ledger, synthesis, and section contracts.
Informal writing	5.6	7.3	Clear but too polished. Needs permission for warmth, contractions, context, and natural rhythm.

6.1 Student Essays

Student essays are the most sensitive use case. The house style can improve student writing by making thesis statements clearer, shortening overloaded sentences, defining terms, and forcing better paragraph discipline. It can also create risk if applied as an authorship detector. Student writing varies by age, assignment, language background, discipline, and teacher expectations. Research on AI detection warns that detectors can be unreliable, adversarially fragile, and biased against non-native English writers.

For student essays, the taxonomy should be converted into a learning rubric rather than an accusation rubric. The strongest mode checks are thesis clarity, claim-evidence-analysis structure, citation integrity, discipline-specific vocabulary, and reflective authorship. A practical essay mode should ask: What is the claim? What evidence supports it? What analysis connects the evidence to the claim? What would this student uniquely know or argue? What sources shaped the answer?

The negative risk is voice flattening. If every student is pushed into executive plain language, essays may lose nuance, exploratory reasoning, and disciplinary style. The neutral view is that sentence clarity and defined terms still help, but the executive guide should not govern essay openings, conclusions, or voice. The positive view is that your system can become a strong anti-bloat tool for student writing if it preserves evidence and analysis as the center of the essay.

6.2 Social Media Posts

Social posts reward compression, clarity, and audience fit. Your current style helps remove hype and filler, which matters because online readers scan and react quickly. Nielsen Norman Group's work on web writing supports concise, scannable, objective content, and your rules map well to that direction.

The gap is that social media is not only a clarity problem. It is also a channel, persona, timing, and interaction problem. A LinkedIn leadership post, a crisis update, a product announcement, and an internal culture post need different levels of warmth, specificity, and polish. The taxonomy is useful here because AI-like posts often sound safe, symmetrical, uplifting, and generic. A social mode should require one clear point, one concrete detail, an audience reason, and a non-generic closing action.

The negative risk is over-polishing. A house style that strips too much personality can make social content feel corporate and synthetic. The neutral view is that the core rules should apply to clarity and accessibility, while hooks, tone, and cadence should remain platform-specific. The positive view is that your style system can become a strong guardrail against hype, vague virtue language, and empty thought leadership.

6.3 Business Documents

Business documents are the strongest domain fit. The executive guide already asks the right questions: what matters, why it matters, what should happen next, what changed, what risks exist, and what decision is required. The taxonomy reinforces this because AI-like business writing often hides behind polished generalities. Your system pushes the writer toward decisions, owners, dates, thresholds, and tradeoffs.

The main enhancement is a one-page decision brief template. It should include decision requested, recommendation, evidence, options considered, tradeoffs, risks, owner, date, and open questions. This is not additional bureaucracy. It is a way to make the existing guide easier to execute.

The negative risk is that business writers may treat the style gate as proof of quality. Passing Vale only means the document avoided specific style problems. It does not prove the recommendation is correct. The positive view is that the current system is already very close to an effective executive communication standard.

6.4 Long-Form Reports

Long-form reports need more than short sentences. They need hierarchy, synthesis, source discipline, section-level purpose, cross-references, visual aids, and a conclusion that does not merely summarize. The uploaded system recognizes some of this through executive structure and repository guidance, but it does not yet define report-level architecture.

The best improvement is a section contract. Each major section should have a purpose, a lead claim, evidence, implications, and a transition. This counters the AI pattern of smooth but shallow coverage. Reports should also include

an evidence ledger for factual claims, a definitions table for technical terms, and a final recommendation or decision path when the report is intended to support action.

The negative risk is under-structuring. A long report that only follows sentence-level rules can still be coherent sentence by sentence and weak as a whole. The neutral view is that some long reports are reference documents, not decision documents, so the executive guide should adapt rather than dominate. The positive view is that your system has the right foundation for report writing because it already values findability and decision usefulness.

6.5 Informal Writing

Informal writing is the weakest fit under the current system. The plain-language rules help, but the executive voice can sound too controlled for messages, quick feedback, team notes, or personal updates. Informal writing often needs warmth, compression, implied context, and a natural rhythm. It may use contractions, sentence fragments, and conversational transitions.

The taxonomy is still useful because AI-like informal writing often becomes strangely polished, overly positive, or emotionally generic. An informal mode should allow contractions, first person, shorter paragraphs, selective fragments, and direct acknowledgment of context. It should still ban unnecessary jargon and vague filler.

The negative risk is authenticity loss. The neutral view is that informal writing should be lightly guided and rarely linted unless it becomes a deliverable. The positive view is that a compact informal mode can make team communication clearer without making it sterile.

7. Recommended Grading Rubric

The following rubric can be used to assess the house style system itself, a domain-specific style mode, or a specific deliverable. It is intentionally practical. It emphasizes outcomes that readers and reviewers can observe.

Dimension	Question	Weight	High-Score Evidence
Reader Fit	Does the document match reader purpose, context, time, and knowledge?	12%	Audience, purpose, channel, and reader action are explicit.
Point and Structure	Does the main point appear early and in the right order?	12%	The lead claim, headings, and flow make the argument easy to follow.
Plain Language	Are sentences, paragraphs, terms, and acronyms clear?	12%	Short direct sentences, defined terms, no avoidable jargon, concise paragraphs.
Evidence Discipline	Are facts, inferences, and recommendations separated?	14%	Claims have support, assumptions are labeled, and open questions are visible.
Domain Fit	Does the piece respect the genre?	10%	Essay, social, business, report, or informal mode is applied appropriately.
AI-Signature Risk	Does the prose avoid generic, formulaic, over-polished patterns?	12%	No empty uplift, repetitive structures, fake specificity, or stock conclusions.
Voice and Authenticity	Does the writing sound authored by a person with judgment?	10%	Specific details, perspective, and appropriate tone are present.
Automation Coverage	Can objective rules be checked repeatably?	8%	Vale catches clear violations without blocking creative or rough work.
Risk Controls	Does the process prevent	6%	Warnings, exceptions,

	misuse, false certainty, and unfair judgments?		attribution, and review responsibilities are clear.
Maintainability	Can the system be kept current without becoming heavy?	6%	Rules are few, understandable, tested, and reviewed periodically.

For a single deliverable, use a simple scoring threshold. A score below 6.0 requires revision before handoff. A score from 6.0 to 7.4 is usable but needs targeted review. A score from 7.5 to 8.9 is strong. A score of 9.0 or higher is ready for high-stakes use after normal factual review. The threshold should be stricter for student integrity reviews, regulated business content, safety communications, and executive decisions.

8. Recommended Improvements

The improvement plan should preserve the system's light operating model. The best next version adds a small number of objective checks, gives authors better examples, and introduces domain modes. It should avoid a large AI-detection layer and should avoid turning vocabulary lists into hard bans.

Priority	Improvement	Impact	Practical Action
1	Add domain modes	High	Create five short add-ons for essays, social posts, business documents, long-form reports, and informal writing.
2	Add sentence and paragraph length warnings	High	Use suggestions only. These rules catch obvious complexity without requiring perfect metrics.
3	Add filler and generic conclusion warnings	High	Flag phrases such as 'it is important to note', 'in conclusion', and empty uplift endings.
4	Add acronym candidate detection	Medium	Warn on repeated all-caps terms. Require manual confirmation of first-use definition.
5	Add an exception policy	Medium	Allow intentional acceptance with a brief reason, especially for source quotes and domain terms.
6	Add before-and-after examples	High	Show how to revise executive prose, essays, social posts, and informal messages.
7	Add evidence handling rules	High	Define source expectations for student essays, reports, and business recommendations.
8	Add script flags	Medium	Use command-line options such as --all, --changed, and --strict instead of environment variables only.
9	Add copied-source exceptions	Medium	Keep vale off blocks or source-pack paths from creating false positives.
10	Review rules quarterly	Low	Keep the jargon list current,

			remove noisy rules, and check whether agents follow the workflow.

9. Proposed Revised House Style Language

The core style guide should stay short. The following revision preserves the current intent while adding evidence, domain fit, and exceptions.

```
# House Style: Plain Language

## Core Rules

1. Write for the reader's purpose, context, and time.
2. State the main point early.
3. Use plain language with short, direct sentences.
4. Keep paragraphs short. Use one main idea per paragraph when practical.
5. Avoid jargon and buzzwords when simple words work.
6. Avoid acronyms. If an acronym is necessary, define it on first use as Full Term (ACR).
7. Use concrete nouns, strong verbs, numbers, dates, owners, and examples when available.
8. Separate facts, assumptions, inferences, recommendations, and open questions.
9. Do not use Unicode em dash characters.
10. Avoid contrast formulas such as 'not X, it is Y'. State the point directly.
11. Use the domain mode when the deliverable is an essay, social post, business document,
long-form report, or informal message.
12. If a warning is intentionally accepted, keep the reason brief and preserve the reader's
need.

## Review Checklist

- Reader and purpose are clear.
- Main point appears early.
- Terms and acronyms are defined or removed.
- Claims have support, or uncertainty is stated.
- Recommendations and next steps are concrete when action is needed.
- Tone fits the domain and channel.
- Vale style gate was run at the checkpoint.
```

10. Proposed Vale Rule Additions

Vale should enforce only rules that are objective enough to check safely. The local rules should remain suggestions or warnings except for the Unicode em dash rule. This keeps the gate useful without making it punitive. Rule files should use the .yml extension because Vale detects rules by that extension.

10.1 Filler Phrases

```
# .vale/styles/PlainEnglish/Filler.yml
extends: existence
message: "Remove filler phrase '%s' unless it is necessary for meaning."
level: suggestion
ignorecase: true
tokens:
  - "it is important to note"
  - "it should be noted"
  - "needless to say"
  - "in conclusion"
  - "in summary"
```

```
  - "at the end of the day"
  - "moving forward"
  - "as we navigate"
```

10.2 Comma Load as a Sentence Complexity Proxy

```
# .vale/styles/PlainEnglish/CommaLoad.yml
extends: occurrence
message: "This sentence has more than 3 commas. Consider splitting it."
level: suggestion
scope: sentence
max: 3
token: ","
```

10.3 Repetitive Stock Conclusions

```
# .vale/styles/PlainEnglish/StockConclusions.yml
extends: existence
message: "This sounds like a stock conclusion. State the decision, implication, or next step
directly."
level: warning
ignorecase: true
tokens:
  - "serves as a testament"
  - "stands as a testament"
  - "beacon of hope"
  - "paving the way"
  - "the future is bright"
  - "only time will tell"
```

10.4 Acronym Candidates

```
# .vale/styles/PlainEnglish/AcronymCandidates.yml
extends: existence
message: "Possible acronym '%s'. Confirm it is defined on first use."
level: suggestion
nonword: true
tokens:
  - "\b[A-Z]{2,6}s?\b"
exceptions:
  - "AI"
  - "IT"
  - "URL"
  - "PDF"
  - "API"
```

10.5 Expanded Substitutions

```
# .vale/styles/PlainEnglish/Jargon.yml
extends: substitution
message: "Replace '%s' with plain language: '%s'."
level: suggestion
ignorecase: true
swap:
  leverage: use
  utilize: use
  synergy: working together
  paradigm: approach
  robust: reliable
```

```
optimize: improve
stakeholder: reader, owner, or affected group
enablement: support
alignment: agreement
actionable: useful or specific
transformative: major or important
impactful: effective
```

These examples should be tested on a representative sample before adoption. The acronym rule will be noisy because it cannot reliably know whether the first use was defined. The jargon list should stay short and should not ban necessary technical terms. The stock conclusion rule should be tuned because some phrases may be legitimate in quoted source material.

11. Risk Analysis and Mitigation

The main risk is misuse. The taxonomy can tempt teams to treat visible writing patterns as proof of AI authorship. That would be unsafe, especially for student essays and non-native English writers. The system should explicitly state that style signals are review prompts, not evidence of misconduct.

Risk	Where It Matters	Mitigation
False accusation risk	Highest in student essays and hiring contexts.	Do not use style signals as proof. Require process evidence, drafts, sources, and conversation.
Voice flattening	High in essays, social posts, and informal writing.	Add domain modes that allow appropriate voice variance.
Overfitting to AI vocabulary	Medium across all domains.	Flag repeated patterns, not isolated words. Treat vocabulary as one clue.
Automation false confidence	High for business documents.	State that Vale passing is a style result, not a quality or factual approval.
Source-pack false positives	Medium for analysis repositories.	Preserve vale off blocks or add source-pack exceptions.
Gate avoidance	Medium because Vale is soft by default.	Use strict mode for final published deliverables and high-stakes content.
Outdated rules	Medium because AI writing changes.	Review rules quarterly and remove noisy checks.

12. Opinionated Perspectives on the Improvement Areas

The following judgments are deliberately opinionated. They separate changes that should be adopted now from changes that should be treated carefully.

12.1 Positive

Keep the executive guide separate. It serves a narrower audience and should not force essays, social posts, or informal writing into a decision-brief structure. The separation is one of the system's best design choices.

Add before-and-after examples. This is the fastest way to make judgment-based rules teachable. A small example bank will help Codex, Claude, and human authors converge on the same style without adding policy weight.

Add sentence and paragraph warnings. They are simple, practical, and aligned with the current house style. Keep them as suggestions so authors can preserve necessary complexity.

Add an exception policy. Exceptions protect judgment. They also prevent the style gate from turning into a mechanical ritual.

12.2 Negative

Do not create an AI-authorship score from the taxonomy. That would overstate the evidence and create avoidable ethical risk. The taxonomy is strongest as a quality lens.

Do not ban AI vocabulary as a hard rule. Words such as robust, significant, crucial, and transformative can be legitimate in context. The problem is repetition, vagueness, and stock phrasing.

Do not force social or informal writing through the executive guide. That would make natural writing feel sterile. Use the core plain-language rules and a lighter channel mode.

12.3 Neutral or Conditional

Acronym enforcement is useful but noisy. Use it as a suggestion and maintain a small accepted acronym list.

Strict mode is appropriate for final external deliverables, regulated content, and executive packages. It is too heavy for routine drafts.

LLM review can help, but only as a second reviewer. Ask the model to identify unclear claims, missing evidence, and generic phrasing. Do not ask it to decide authorship.

13. Recommended Operating Model

The daily model should be easy to explain and easy to run. The author drafts naturally, chooses the right domain mode, edits for substance, applies the house style, runs the style gate at the checkpoint, fixes the flagged text, and records any intentional exceptions. For executive documents, the author also confirms that the recommendation, tradeoffs, risks, and next steps are visible in under one minute.

Role	Expected Behavior
Human authors	Use the domain checklist before final edit. Run Vale at handoff. Record exceptions in the document or review notes.
Codex	Follow repository rules, gate only deliverables, preserve copied-source blocks, and avoid changing meaning when fixing lint.
Claude	Apply the same house style and domain modes. Use the style gate when working in the repository.
Reviewers	Check substance first, then style. Do not treat a clean Vale result as approval of facts or recommendations.

14. Conclusion

Your house style system is already strong where it matters most for business communication: it makes writing clear, direct, and decision-useful without overloading the workflow. The AI-writing detection taxonomy exposes the next maturity step. Strong writing is not only plain. It is specific, situated, evidence-aware, appropriately voiced, and resistant to formulaic patterns.

The practical path forward is incremental. Keep the core house style short. Keep the executive guide separate. Add domain modes. Expand Vale only where checks are objective. Add examples and a small exception policy. Treat AI-detection clues as quality risks, not proof of authorship. With those changes, the system can move from a good plain-language gate to a broader writing-quality operating system that remains practical for daily use.

Appendix A. Domain Mode Cards

Mode	Checklist
Student Essay Mode	Thesis stated; claims supported; evidence cited; analysis connects evidence to claim; student voice preserved; uncertainty acknowledged.
Social Post Mode	One point; specific audience; concrete detail; platform-appropriate tone; useful call to action; no empty uplift.
Business Document Mode	Decision or update stated early; recommendation clear; risks and tradeoffs explicit; owner and date visible; next step concrete.
Long-Form Report Mode	Executive summary; section contracts; evidence ledger; definitions; synthesis; action path or research conclusion.
Informal Mode	Context first; natural tone; useful brevity; contractions allowed; empathy or warmth when appropriate; no corporate filler.

References

Department of Labor. Plain Language Quick Reference Guide. https://www.dol.gov/sites/dolgov/files/general/Plain-Language-Quick-Reference-Guide.pdf

Digital.gov. Plain Language Guide Series. https://digital.gov/guides/plain-language

Dugan, L., Hwang, A., Trhlik, F., Zhu, A., Ludan, J. M., Xu, H., Ippolito, D., and Callison-Burch, C. 2024. RAID: A Shared Benchmark for Robust Evaluation of Machine-Generated Text Detectors. Proceedings of ACL 2024. https://aclanthology.org/2024.acl-long.674/

ISO. 2023. ISO 24495-1:2023 Plain language - Part 1: Governing principles and guidelines. Preview sample. https://cdn.standards.iteh.ai/samples/78907/d194fac21d6a45f38bfcfec9657f7498/ISO-24495-1-2023.pdf

Microsoft. Microsoft Writing Style Guide. https://learn.microsoft.com/en-us/style-guide/welcome/

Microsoft. Writing tips - Microsoft Style Guide. https://learn.microsoft.com/en-us/style-guide/global-communications/writing-tips

Morkes, J., and Nielsen, J. 1997. Concise, Scannable, and Objective: How to Write for the Web. Nielsen Norman Group. https://www.nngroup.com/articles/concise-scannable-and-objective-how-to-write-for-the-web/

Russell, J., Karpinska, M., and Iyyer, M. 2025. People who frequently use ChatGPT for writing tasks are accurate and robust detectors of AI-generated text. Proceedings of ACL 2025. https://github.com/jenna-russell/human_detectors

Sadasivan, V. S., Kumar, A., Balasubramanian, S., Wang, W., and Feizi, S. 2024. Can AI-Generated Text be Reliably Detected? ICLR 2024. https://arxiv.org/abs/2303.11156

Stanford HAI. 2023. AI-Detectors Biased Against Non-Native English Writers. https://hai.stanford.edu/news/ai-detectors-biased-against-non-native-english-writers

Vale CLI. Styles and rule documentation. https://vale.sh/docs/styles

Vale CLI. Occurrence rule documentation. https://vale.sh/docs/checks/occurrence

Blog-context sources

Fast Company. 2026. AI power users are pulling away from everyone else, Microsoft says. Yahoo/Tech syndicated copy of Fast Company coverage.

Microsoft. 2026. 2026 Work Trend Index Annual Report: Agents, human agency, and the opportunity for every organization.